



OBSERRA

Executive Protection & Intelligence LLC

Obserra Agentic AI Security

Admin & Operator Guide

Agentic AI Security Control Plane — machine authority, guardrails & runtime enforcement

August 2026

Confidential — prepared for internal use.

Contents

1. About Obserra Agentic AI Security
2. Installing the App (One-Click PWA)
3. On-Premise Installation (Docker)
4. Signing In
5. Executive Overview
6. Agentic AI Security Control Plane
7. Agent Inventory
8. Tool Toxicity Map
9. Guardrails & Red Team
10. Shadow AI Discovery
11. Kill Switch & Runtime Enforcement
12. AI Incidents & Workflows
13. Defensibility & Evidence
14. Board Brief Scheduler
15. Settings, Branding & Deployment
16. Obserrian Advisor
17. Support

1. About Obserra Agentic AI Security

Obserra Agentic AI Security is the control plane for the AI agents and models operating across your enterprise. It gives security, GRC and executive teams a single live view of which agents exist, what tools and permissions they hold, how much autonomy they have, and where their capabilities combine into dangerous — 'toxic' — patterns.

Every number is computed LIVE from the underlying agent records on each request (No-Mock): delegated authority tiers, modelled agent risk scores, guardrail coverage and the heuristic red-team baseline are all derived from real telemetry. Governance actions (sanction, restrict, suspend, kill) are written to the Defensibility Ledger so every decision is audit-defensible, and the runtime enforcement connector can push those decisions to an external agent runtime.

2. Installing the App (One-Click PWA)

Obserra installs like a native app straight from the browser — no app store required — and works across desktop, tablet and mobile.

Desktop (Chrome / Edge): click the Install icon in the address bar, or use the in-app 'Install' banner. Android (Chrome): tap the 'Install' banner or menu -> Add to Home screen. iPhone / iPad (Safari): tap Share -> Add to Home Screen.

Once installed, the app launches full-screen and can receive push notifications for AI security alerts, shadow-AI discoveries and guardrail breaches.

3. On-Premise Installation (Docker)

For fully self-hosted deployments, download the on-premise package from Settings -> Deployment & Documentation. An optional `install.sh` is included.

Prerequisites: Docker 24+ and Docker Compose v2, 2 vCPU / 4 GB RAM (8 GB recommended).

Steps: (1) place `backend/` and `frontend/` next to the `deploy/` folder; (2) copy `.env.example` to `.env` and set `JWT_SECRET`, `EMERGENT_LLM_KEY` and `PUBLIC_URL`; (3) run `docker compose -f deploy/docker-compose.yml --env-file deploy/.env up -d --build`; (4) open `http://:8080`. Full details are in the bundled `INSTALL.md`.

4. Signing In

Open the app and sign in with your work email and password. Passwords follow NIST 800-63B (≥ 12 chars with upper/lower/number/symbol). Roles determine what you see: admins get full enforcement controls, executives get the board view.



Figure: Signing In

5. Executive Overview

The landing dashboard is the complete rollup of the AI security estate: modelled agent risk, autonomous agents, toxic capability combinations, shadow-AI exposure, guardrail gaps and open AI incidents — all as one board-ready view. Every KPI and card opens a standardized deep-dive with an AI strategic brief and recommended actions, and admins can email the executive brief on demand or on a cadence from here.




Figure: Executive Overview

6. Agentic AI Security Control Plane

The dedicated seven-tab workspace: Mission Control, Agent Inventory, Authority & Tools (with the Tool Toxicity Map), Guardrails & Red Team, Shadow AI, Incidents and Defensibility. Mission Control summarises the estate; every tile is clickable and drills into the relevant tab or agent.

7. Agent Inventory

Every registered AI agent with its owner, model, delegated tools and permissions, modelled risk score, authority tier, guardrail coverage and governance status. Search and filter by status or risk class; open any agent for full evidence and AI analysis.

8. Tool Toxicity Map

A visual Agent -> Tool -> Permission -> Resource graph that flags toxic capability combinations — for example an action-capable tool such as shell.exec paired with write permissions and no human-in-the-loop guardrail, or a data-exfiltration-capable tool with access to sensitive resources. The heatmap makes the most dangerous agents glanceable so reviewers can neutralise them first.

9. Guardrails & Red Team

Enterprise guardrail coverage (input/output filtering, tool allowlist, human approval) across every agent, plus the heuristic red-team baseline. Admins can toggle a guardrail on an agent's governance record and run the deterministic red-team probes; results are explicitly labelled as a heuristic baseline, not live adversarial runtime testing.

10. Shadow AI Discovery

A discovery feed that auto-populates the review queue with unsanctioned AI systems detected across the estate — common public GenAI SaaS as well as agents flagged shadow. Admins run discovery in one click and sanction each system to bring it under governance.

11. Kill Switch & Runtime Enforcement

From any agent's detail view, admins can Suspend (restrict), Kill (block) or Resume an agent. The runtime enforcement connector flips the agent's governance status, records the action to the Defensibility Ledger, posts a Slack/Teams alert, and — when an agent-runtime webhook is configured — dispatches the enforcement command to the external execution environment. The UI honestly reflects whether enforcement was applied in the control plane only or pushed to an external runtime.

12. AI Incidents & Workflows

AI security incident records (severity, mode, status) from the governance backend, alongside the related governance workflows. Each incident opens a deep-dive with an AI brief and recommended actions.

13. Defensibility & Evidence

The evidence layer: live data-source status, an explicit separation of FACT vs MODELLED vs HEURISTIC BASELINE vs AI RECOMMENDATION, the runtime-enforcement boundary statement and connector health. This is what makes the whole plane audit-defensible.

14. Board Brief Scheduler

A one-click control (admin) that emails the AI Security Executive Brief — modelled agent risk, autonomous agents, toxic combinations, shadow AI and open incidents — via the managed email pipeline. Choose a weekly or monthly cadence, or send it immediately.

15. Settings, Branding & Deployment

Personal preferences (digest cadence, replay the guided tour) plus admin controls: alert recipients, custom branding (company name, logo, accent colour), a 'send me a test now' button, and the Deployment & Documentation downloads (on-premise package and these guides in PDF and Word).




Figure: Settings, Branding & Deployment

16. Obserrian Advisor

The floating Obserrian Advisor (bottom-right, and via the top-bar ask box) answers AI security questions grounded on your LIVE agent telemetry — delegated tools, permissions, guardrail coverage, tool-governance violations and the red-team baseline. It can execute recommended actions and, for admins, reports its own usage and spend.

17. Support

For assistance contact your Obserra administrator. Risk scores and AI evaluations are decision-support estimates and do not constitute legal, financial, regulatory or security guarantees.